

Machine Learning Exercise sheet 09
Deep Learning

Anastasia Stamatouli : 03710902

15/12/2019

Problem (1) :

$$y = \log \sum_{i=1}^n e^{x_i} \Leftrightarrow$$

$$e^y = \sum_{i=1}^n e^{x_i} \Leftrightarrow e^{-a} e^y = e^{-a} \sum_{i=1}^n e^{x_i} \Leftrightarrow$$

$$e^{y-a} = \sum_{i=1}^n e^{-a} e^{x_i} \Leftrightarrow e^{y-a} = \sum_{i=1}^n e^{x_i-a} \Leftrightarrow$$

$$y-a = \log \sum_{i=1}^n e^{x_i-a} \Leftrightarrow \boxed{y = a + \log \sum_{i=1}^n e^{x_i-a}}$$

Problem (2) :

$$\frac{e^{-a}}{e^{-a}} \frac{e^{x_i}}{\sum_{i=1}^N e^{x_i}} = \frac{e^{x_i-a}}{\sum_{i=1}^N e^{x_i-a}} \quad \text{for any constant } a.$$

exercise_09_notebook

December 15, 2019

```
[1]: import numpy as np
import matplotlib.pyplot as plt

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import label_binarize
from sklearn.metrics import accuracy_score

from scipy.special import softmax

[2]: X, y = load_digits(return_X_y=True)
# Convert a categorical vector y (shape [N]) into a one-hot encoded matrix
# → (shape [N, K])
Y = label_binarize(y, np.unique(y)).astype(np.float64)

np.random.seed(123)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25)

[3]: N, K = Y.shape # N - num_samples, K - num_classes
D = X.shape[1] # num_features
```

Remember from the tutorial: 1. No for loops! Use matrix multiplication and broadcasting whenever possible. 2. Think about numerical stability

```
[4]: import nn_utils # module containing helper functions for checking the
# → correctness of your code
```

0.1 Task 1: Affine layer

Implement forward and backward functions for Affine layer

```
[5]: class Affine:
    def forward(self, inputs, weight, bias):
        """Forward pass of an affine (fully connected) layer.

        Args:
            inputs: input matrix, shape (N, D)
            weight: weight matrix, shape (D, H)
```

```

        bias: bias vector, shape (H)

Returns
        out: output matrix, shape (N, H)
        """
        self.cache = (inputs, weight, bias)
        #####
        #

        out = np.dot(inputs, weight) + bias

        #####
        assert out.shape[0] == inputs.shape[0]
        assert out.shape[1] == weight.shape[1] == bias.shape[0]
        return out

def backward(self, d_out):
    """Backward pass of an affine (fully connected) layer.

    Args:
        d_out: incoming derivatives, shape (N, H)

    Returns:
        d_inputs: gradient w.r.t. the inputs, shape (N, D)
        d_weight: gradient w.r.t. the weight, shape (D, H)
        d_bias: gradient w.r.t. the bias, shape (H)
    """
    inputs, weight, bias = self.cache
    #####
    #

    d_inputs = np.dot(d_out, weight.T)
    d_weight = np.dot(inputs.T, d_out)
    d_bias = np.sum(d_out, axis=0)

    #####
    assert np.all(d_inputs.shape == inputs.shape)
    assert np.all(d_weight.shape == weight.shape)
    assert np.all(d_bias.shape == bias.shape)
    return d_inputs, d_weight, d_bias

```

```

[6]: affine = Affine()
      nn_utils.check_affine(affine)

```

All checks passed succesfully!

0.2 Task 2: ReLU layer

Implement forward and backward functions for ReLU layer

```
[7]: class ReLU:
    def forward(self, inputs):
        """Forward pass of a ReLU layer.

        Args:
            inputs: input matrix, arbitrary shape

        Returns:
            out: output matrix, has same shape as inputs
        """
        self.cache = inputs
        #####
        #
        out = np.maximum(inputs, 0)

        #####
        assert np.all(out.shape == inputs.shape)
        return out

    def backward(self, d_out):
        """Backward pass of an ReLU layer.

        Args:
            d_out: incoming derivatives, same shape as inputs in forward

        Returns:
            d_inputs: gradient w.r.t. the inputs, same shape as d_out
        """
        inputs = self.cache
        #####
        #
        d_inputs = np.array(d_out, copy=True)
        d_inputs[inputs <= 0] = 0
        d_inputs[inputs > 0] = d_out [inputs > 0]

        #####
        assert np.all(d_inputs.shape == inputs.shape)
        return d_inputs
```

```
[8]: relu = ReLU()
nn_utils.check_relu(relu)
```

All checks passed succesfully!

0.3 Task 3: CategoricalCrossEntropy layer

Implement forward and backward for CategoricalCrossEntropy layer

```
[9]: class CategoricalCrossEntropy:
    def forward(self, logits, labels):
        """Compute categorical cross-entropy loss.

        Args:
            logits: class logits, shape (N, K)
            labels: target labels in one-hot format, shape (N, K)

        Returns:
            loss: loss value, float (a single number)
        """
        #####
        #

        num_examples = labels.shape[0]
        # get unnormalized probabilities

        labels2=labels.argmax(axis=1)

        exp_scores = np.exp(logits)
        # normalize them for each example
        probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)

        correct_logprobs = -np.log(probs[range(num_examples), labels2])

        loss = np.sum(correct_logprobs)/num_examples

        #####
        # probs is the (N, K) matrix of class probabilities

        self.cache = (probs , labels)
        assert isinstance(loss, float)
        return loss

    def backward(self, d_out=1.0):
        """Backward pass of the Cross Entropy loss.

        Args:
            d_out: Incoming derivatives. We set this value to 1.0 by default,
                    since this is the terminal node of our computational graph
                    (i.e. we usually want to compute gradients of loss w.r.t.
```

```

        other model parameters).

Returns:
    d_logits: gradient w.r.t. the logits, shape (N, K)
    d_labels: gradient w.r.t. the labels
               we don't need d_labels for our models, so we don't
               compute it and set it to None. It's only included in the
               function definition for consistency with other layers.
"""
probs, labels = self.cache
#####
#
num_examples = labels.shape[0]
labels2=labels.argmax(axis=1)

d_logits = probs
d_logits[range(num_examples),labels2] -= 1
d_logits /= num_examples
d_logits *= d_out
#####
d_labels = None
assert np.all(d_logits.shape == probs.shape == labels.shape)
return d_logits, d_labels

```

```

[10]: cross_entropy = CategoricalCrossEntropy()
      nn_utils.check_cross_entropy(cross_entropy)

```

All checks passed successfully!

1 Logistic regression (with backpropagation) — nothing to do in this section

```

[11]: class LogisticRegression:
      def __init__(self, num_features, num_classes, learning_rate=1e-2):
          """Logistic regression model.
             Gradients are computed with backpropagation.

             The model consists of the following sequence of operations:

             input -> affine -> softmax
             """
          self.learning_rate = learning_rate

          # Initialize the model parameters
          self.params = {

```

```

        'W': np.zeros([num_features, num_classes]),
        'b': np.zeros([num_classes])
    }

    # Define layers
    self.affine = Affine()
    self.cross_entropy = CategoricalCrossEntropy()

def predict(self, X):
    """Generate predictions for one minibatch.

    Args:
        X: data matrix, shape (N, D)

    Returns:
        Y_pred: predicted class probabilities, shape (N, D)
        Y_pred[n, k] = probability that sample n belongs to class k
    """
    logits = self.affine.forward(X, self.params['W'], self.params['b'])
    Y_pred = softmax(logits, axis=1)
    return Y_pred

def step(self, X, Y):
    """Perform one step of gradient descent on the minibatch of data.

    1. Compute the cross-entropy loss for given (X, Y).
    2. Compute the gradients of the loss w.r.t. model parameters.
    3. Update the model parameters using the gradients.

    Args:
        X: data matrix, shape (N, D)
        Y: target labels in one-hot format, shape (N, K)

    Returns:
        loss: loss for (X, Y), float, (a single number)
    """
    # Forward pass - compute the loss on training data
    logits = self.affine.forward(X, self.params['W'], self.params['b'])
    loss = self.cross_entropy.forward(logits, Y)

    # Backward pass - compute the gradients of loss w.r.t. all the model
    ↪ parameters
    grads = {}
    d_logits, _ = self.cross_entropy.backward()
    _, grads['W'], grads['b'] = self.affine.backward(d_logits)

    # Apply the gradients

```



```

        for p in self.params:
            self.params[p] = self.params[p] - self.learning_rate * grads[p]
        return loss

```

```

[12]: # Specify optimization parameters
learning_rate = 1e-2
max_epochs = 501
report_frequency = 50

```

```

[13]: log_reg = LogisticRegression(num_features=D, num_classes=K)

```

```

[14]: for epoch in range(max_epochs):
        loss = log_reg.step(X_train, Y_train)
        if epoch % report_frequency == 0:
            print(f'Epoch {epoch:4d}, loss = {loss:.4f}')

```

```

Epoch    0, loss = 2.3026
Epoch   50, loss = 0.2275
Epoch  100, loss = 0.1599
Epoch  150, loss = 0.1306
Epoch  200, loss = 0.1130
Epoch  250, loss = 0.1009
Epoch  300, loss = 0.0918
Epoch  350, loss = 0.0846
Epoch  400, loss = 0.0788
Epoch  450, loss = 0.0738
Epoch  500, loss = 0.0696

```

```

[15]: y_test_pred = log_reg.predict(X_test).argmax(1)
y_test_true = Y_test.argmax(1)

```

```

[16]: print(f'test set accuracy = {accuracy_score(y_test_true, y_test_pred):.3f}')

```

```

test set accuracy = 0.953

```

2 Feed-forward neural network (with backpropagation)

```

[17]: def xavier_init(shape):
        """Initialize a weight matrix according to Xavier initialization.

        See pytorch.org/docs/stable/nn.init#torch.nn.init.xavier\_uniform\_for\_
        ↪ details.
        """
        a = np.sqrt(6.0 / float(np.sum(shape)))
        return np.random.uniform(low=-a, high=a, size=shape)

```

2.1 Task 4: Implement a two-layer FeedForwardNeuralNet model

You can use the LogisticRegression class for reference

```
[18]: class FeedforwardNeuralNet:
    def __init__(self, input_size, hidden_size, output_size, learning_rate=1e-2):
        """A two-layer feedforward neural network with ReLU activations.

        (input_layer -> hidden_layer -> output_layer)

        The model consists of the following sequence of operations:

        input -> affine -> relu -> affine -> softmax

        """
        self.learning_rate = learning_rate

        # Initialize the model parameters
        self.params = {
            'W1': xavier_init([input_size, hidden_size]),
            'b1': np.zeros([hidden_size]),
            'W2': xavier_init([hidden_size, output_size]),
            'b2': np.zeros([output_size]),
        }

        # Define layers
        #####
        self.affine1 = Affine()
        self.relu = ReLU()
        self.affine2 = Affine()
        self.cross_entropy = CategoricalCrossEntropy()

        #####

    def predict(self, X):
        """Generate predictions for one minibatch.

        Args:
            X: data matrix, shape (N, D)

        Returns:
            Y_pred: predicted class probabilities, shape (N, D)
            Y_pred[n, k] = probability that sample n belongs to class k
        """
        #####
        #
```

```

        affloutput = self.affine1.forward(X, self.params['W1'], self.
↪params['b1'])
        reluoutput = self.relu.forward(affloutput)
        logits = self.affine2.forward(reluoutput, self.params['W2'], self.
↪params['b2'])
        Y_pred = softmax(logits, axis=1)

#####
return Y_pred

def step(self, X, Y):
    """Perform one step of gradient descent on the minibatch of data.

    1. Compute the cross-entropy loss for given (X, Y).
    2. Compute the gradients of the loss w.r.t. model parameters.
    3. Update the model parameters using the gradients.

    Args:
        X: data matrix, shape (N, D)
        Y: target labels in one-hot format, shape (N, K)

    Returns:
        loss: loss for (X, Y), float, (a single number)
    """
    #####
    #
    # Forward pass - compute the loss on training data
    affloutput = self.affine1.forward(X, self.params['W1'], self.
↪params['b1'])
    #
    print(self.params['W1'].shape)
    #
    print(self.params['b1'].shape)
    reluoutput = self.relu.forward(affloutput)
    logits = self.affine2.forward(reluoutput, self.params['W2'], self.
↪params['b2'])

    loss = self.cross_entropy.forward(logits, Y)

    # Backward pass - compute the gradients of loss w.r.t. all the model_
↪parameters
    grads = {}
    d_logits, _ = self.cross_entropy.backward()

    d_inputs2, grads['W2'], grads['b2'] = self.affine2.backward(d_logits)

```

```

        d_inputs = self.relu.backward(d_inputs2)

#         print(d_inputs.shape)

        _, grads['W1'], grads['b1'] = self.affine1.backward(d_inputs)

# Apply the gradients
for p in self.params:
    self.params[p] = self.params[p] - self.learning_rate * grads[p]

#####
return loss

```

```
[19]: H = 32 # size of the hidden layer
```

```

# Specify optimization parameters
learning_rate = 1e-2
max_epochs = 501
report_frequency = 50

```

```
[20]: model = FeedforwardNeuralNet(input_size=D, hidden_size=H, output_size=K,
    ↪ learning_rate=learning_rate)
```

```
[21]: for epoch in range(max_epochs):
        loss = model.step(X_train, Y_train)
        if epoch % report_frequency == 0:
            print(f'Epoch {epoch:4d}, loss = {loss:.4f}')
```

```

Epoch    0, loss = 8.5876
Epoch   50, loss = 0.6002
Epoch  100, loss = 0.3517
Epoch  150, loss = 0.2510
Epoch  200, loss = 0.1975
Epoch  250, loss = 0.1631
Epoch  300, loss = 0.1401
Epoch  350, loss = 0.1231
Epoch  400, loss = 0.1098
Epoch  450, loss = 0.0989
Epoch  500, loss = 0.0897

```

```
[22]: y_test_pred = model.predict(X_test).argmax(1)
        y_test_true = Y_test.argmax(1)
```

```
[23]: print(f'test set accuracy = {accuracy_score(y_test_true, y_test_pred):.3f}')
```

```
test set accuracy = 0.938
```

```
[ ]:
```